

Computational Communication Science 2

Week 2 - Lecture

Bottom-up approaches to text analysis: From preprocessing to vectorization

Anne Kroon

a.c.kroon@uva.nl, @annekroon

April 7, 2025

Digital Society Minor, University of Amsterdam

Today's Agenda

Recap

Building text representations: vectorizers

General idea

Pruning

Introducing embeddings

Word embeddings

Recap

Review of key concepts

- **Text preprocessing:** Cleaning and preparing raw text for analysis.
- **Core techniques:**
 - tokenization
 - Stopword removal
 - Using built-in string methods for text cleaning
- **Methodological approaches:** Combining *bottom-up* (data-driven) and *top-down* (theory-driven) approaches.

Typical preprocessing steps

Preprocessing steps

tokenization How do we (best) split a sentence into tokens (terms, words)?

pruning How can we remove unnecessary words/punctuation?

lemmatization and stemming How can we make sure that slight variations of the same word are not counted differently?

ngrams Neighbouring terms

Simple Tokenization using `.split()`

Tokenization: Splitting text into words or subwords is essential for many NLP tasks.

Simple Tokenization with `.split()`: Python's built-in `'split()'` method is a simple way to break a text into tokens by spaces.

```
1 text = "This is an example sentence."  
2 tokens = text.split()  
3 print(tokens) # Output: ['This', 'is', 'an', 'example', 'sentence.']
```

Advanced Tokenization with TreebankWordTokenizer

For more sophisticated tokenization that handles punctuation, contractions, etc., we use 'TreebankWordTokenizer' from NLTK.

```
1 from nltk.tokenize import TreebankWordTokenizer
2
3 docs = ["This is an example sentence."]
4 tokens = [TreebankWordTokenizer().tokenize(d) for d in docs]
5 print(tokens) # Output: [['This', 'is', 'an', 'example', 'sentence',
  ↳  '.']]
```

Understanding N-grams

- An n-gram is a sequence of n words treated as a single feature.
- Examples:
 - Unigrams (1-word units): "science"
 - Bigrams (2-word units): "data science"
 - Trigrams (3-word units): "machine learning model"
- **Why use n-grams?** Captures context and word relationships beyond single words.
- **Collocations:** Identifying n-grams that occur frequently and naturally together (e.g., "data science," "machine learning") helps uncover meaningful word patterns in text.

Generating N-grams with Python

Code Example: Generating n-grams using Python's NLTK library.

```
1 import nltk
2 from nltk.util import ngrams
3 from nltk.tokenize import word_tokenize
4
5 # Sample text
6 text = "Machine learning is a powerful tool for data science."
7
8 # Tokenize the text
9 tokens = word_tokenize(text)
10
11 # Create bigrams and trigrams
12 bigrams = list(ngrams(tokens, 2))
13 trigrams = list(ngrams(tokens, 3))
14
15 # Count the frequency of n-grams
16 print("Bigrams:", bigrams)
17 print("Trigrams:", trigrams)
```

Identifying collocations

Code Example: Finding collocations using NLTK's BigramCollocationFinder.

```
1  from nltk.collocations import BigramCollocationFinder,  
   ↪ BigramAssocMeasures  
2  from collections import Counter  
3  
4  # Find bigram collocations using NLTK's BigramCollocationFinder  
5  bigram_finder = BigramCollocationFinder.from_words(tokens)  
6  bigram_collocations =  
   ↪ bigram_finder.nbest(BigramAssocMeasures.likelihood_ratio, 5)  
7  
8  print("Collocations (Bigram):", bigram_collocations)  
9
```

Stemming: Reducing Words to Their Root

- **Stemming** chops off word endings to reduce words to a common root.
- It's fast and simple, but may produce non-dictionary words.
- Example: "running" → "run", "flies" → "fli"

Code example (PorterStemmer):

```
1 from nltk.stem import PorterStemmer
2
3 stemmer = PorterStemmer()
4 words = ["running", "flies", "easily", "fairly"]
5 stems = [stemmer.stem(word) for word in words]
6
7 print(stems)
8 # Output: ['run', 'fli', 'easili', 'fairli']
```

Lemmatization: Getting the Dictionary Form

- **Lemmatization** maps words to their base (dictionary) form.
- More accurate than stemming, but slower and requires POS tagging for best results.
- Example: “better” → “good”, “running” → “run”

Code Example (WordNetLemmatizer):

```
1 from nltk.stem import WordNetLemmatizer
2 from nltk.corpus import wordnet
3 from nltk import pos_tag, word_tokenize
4
5 lemmatizer = WordNetLemmatizer()
6 words = ["running", "flies", "better"]
7 lemmas = [lemmatizer.lemmatize(word) for word in words]
8
9 print(lemmas)
10 # Output: ['running', 'flies', 'better']
```

Building text representations: vectorizers

Building text representations: vectorizers

General idea

A text as a collections of word

Let us represent a string

```
1 t = "This this is is is a test test test"
2 # like this:
3 print(Counter(t.split()))
```

```
1 Counter({'is': 3, 'test': 3, 'This': 1, 'this': 1, 'a': 1})
```

Compared to the original string, this representation

- is less repetitive
- preserves word frequencies
- but does *not* preserve word order
- can be interpreted as a vector to calculate with (!!!)

Of course, still a lot of stuff to fine-tune. . . (for example, This/this)

From vector to matrix

If we do this for multiple texts, we can arrange the vectors in a table.

`t1 = "This this is is is a test test test"`

`t2 = "This is an example"`

	a	an	example	is	this	This	test
<i>t1</i>	1	0	0	3	1	1	3
<i>t2</i>	0	1	1	1	0	1	0



What can you do with such a matrix? Why would you want to represent a collection of texts in such a way?

What is a vectorizer

- Transforms a list of texts into a sparse (!) matrix (of word frequencies)
- Vectorizer needs to be “fitted” to the training data (learn which words (features) exist in the dataset and assign them to columns in the matrix)
- Vectorizer can then be re-used to transform other datasets

The cell entries: raw counts versus tf-idf scores

- In the example, we entered simple counts (the “term frequency”)



But are all terms equally important?

The cell entries: raw counts versus tf-idf scores

- In the example, we entered simple counts (the “term frequency”)
- But does a word that occurs in almost all documents contain much information?
- And isn’t the presence of a word that occurs in very few documents a pretty strong hint?
- **Solution:** Weigh by *the number of documents in which the term occurs at least once* (the “document frequency”)

⇒ we multiply the “term frequency” (tf) by the inverse document frequency (idf)

(usually with some additional logarithmic transformation and normalization applied, see https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.TfidfTransformer.html)

TF-IDF: Weighted Importance of Words

$$\text{tf-idf} = \text{tf}_{i,j} \times \log \left(\frac{N}{\text{df}_i} \right)$$

- $\text{tf}_{i,j}$ = term frequency of term i in document j
- df_i = number of documents containing the term i
- N = total number of documents in the corpus
- **Term Frequency (TF):** Measures how often a term appears in a document. More frequent terms in a document receive a higher TF score.
- **Inverse Document Frequency (IDF):** Downscales terms that appear in many documents and boosts those that are rare but important.

Is tf-idf always better?

It depends.

- Ultimately, it's an empirical question which works better (→ machine learning)
- In many scenarios, “discounting” too frequent words and “boosting” rare words makes a lot of sense (most frequent words in a text can be highly un-informative)
- Beauty of raw tf counts, though: interpretability + describes document in itself, not in relation to other documents

Different vectorizers

1. CountVectorizer (=simple word counts)
2. TfidfVectorizer (word counts (“term frequency”) weighted by number of documents in which the word occurs at all (“inverse document frequency”))

Internal representations

Sparse vs dense matrices

- → tens of thousands of columns (terms), and one row per document
- Filling all cells is inefficient *and* can make the matrix too large to fit in memory (!!!)
- Solution: store only non-zero values with their coordinates! (sparse matrix)
- dense matrix (or dataframes) not advisable, only for toy examples

s p a r s e

0	7	0	0	0	0	6
0	7	6	3	0	4	0
0	4	3	0	0	0	0
4	2	0	0	0	0	0
0	0	0	0	3	2	4

© Matt Eding

DENSE

0	7	0	0	0	0	6
0	7	6	3	0	4	0
0	4	3	0	0	0	0
4	2	0	0	0	0	0
0	0	0	0	3	2	4

<https://matteding.github.io/2019/04/25/sparse-matrices/>

We learned in week 1 how to tokenize with a list comprehension (and that's often a good idea!).

```
1 from nltk.tokenize import TreebankWordTokenizer
2 tokens = [TreebankWordTokenizer().tokenize(d) for d in docs]
```

But what if we want to *directly* get a DTM instead of lists of tokens?

OK, good enough, perfect?

scikit-learn's CountVectorizer (default settings)

- applies lowercasing
- deals with punctuation etc. itself
- minimum word length > 1
- more technically, tokenizes using this regular expression:
`r"(?u)\b\w\w+\b"`¹

```
1 from sklearn.feature_extraction.text import CountVectorizer
2 cv = CountVectorizer()
3 dtm_sparse = cv.fit_transform(docs)
```

¹?u = support unicode, \b = word boundary

OK, good enough, perfect?

CountVectorizer supports more

- stopword removal
- custom regular expression
- or even using an external tokenizer
- ngrams instead of unigrams

see

https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html

Best of both worlds

Use the Count vectorizer with a NLTK-based external tokenizer! (see book)

This notebook might help to better your understanding of vectorizers!

Building text representations: vectorizers

Pruning

General idea

- Idea behind both stopwords removal and tf-idf: too frequent words are uninformative
- (possible) downside stopwords removal: a priori list, does not take empirical frequencies in dataset into account
- (possible) downside tf-idf: does not reduce number of features

Pruning: remove all features (tokens) that occur in less than X or more than X of the documents

CountVectorizer, only stopwords removal

```
1 from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
2 myvectorizer = CountVectorizer(stop_words=mystopwords)
```

CountVectorizer, better tokenization, stopwords removal (pay attention that stopwords list uses same tokenization!):

```
1 myvectorizer = CountVectorizer(tokenizer = TreebankWordTokenizer().tokenize,
  ↳ stop_words=mystopwords)
```

Additionally remove words that occur in more than 75% or less than $n = 2$ documents:

```
1 myvectorizer = CountVectorizer(tokenizer = TreebankWordTokenizer().tokenize,
  ↳ stop_words=mystopwords, max_df=.75, min_df=2)
```

All together: tf-idf, explicit stopwords removal, pruning

```
1 myvectorizer = TfidfVectorizer(tokenizer = TreebankWordTokenizer().tokenize,
  ↳ stop_words=mystopwords, max_df=.75, min_df=2)
```




*What is “best”? Which
(combination of) techniques to
use, and how to decide?*

Exercise time: Word cloud

Let's put this into practice!

- Follow the instructions in the exercise material.
- You will create a word cloud from text data.
- Apply preprocessing techniques such as tokenization, stopwords removal, and normalization.

Exercise Link: [GitHub: Word cloud exercise](#) *Think about how*

different preprocessing choices impact the resulting word cloud

From text to features

- Text needs to be transformed into numerical representations for computational analysis.
- Tokenization breaks text into meaningful units (words, phrases, n-grams).
- Frequency-based representations allow us to quantify text characteristics.

Introducing embeddings


Why move beyond TF-IDF?

Limitations of TF-IDF and CountVectorizer:

- Create **sparse**, high-dimensional representations.
- Ignore **semantic meaning**—similar words have different vectors.
- Cannot capture **contextual relationships**.

Solution: Word embeddings


Embeddings create **dense**, low-dimensional representations that retain meaning and context.



What are embeddings?

- Map words or documents into **continuous vector spaces**.
- Words with similar meanings have **similar vectors**.
- **Learned** from large text corpora using machine learning models.

Popular Examples: Word2Vec, GloVe, fastText, BERT embeddings.



Introducing embeddings


Word embeddings


Understanding embeddings


What are word embeddings?

- No technical details here, just the general idea
- Word embeddings help capture the meaning of text
- Word embeddings are low-dimensional vector representations that capture semantic meaning
- Used to be state-of-the-art in NLP (but now: contextualized embeddings, e.g., BERT or GPT)
- *“...a word is characterized by the company it keeps...”* (Firth, 1957)











How embeddings work

- Each word (or document) is represented by a **vector** in a high-dimensional space.
- The model learns these vectors by predicting word co-occurrences in text.
- Example: **Word2Vec** uses a neural network to predict surrounding words.

Example: Word similarity with embeddings

Word2Vec analogy:

- $\text{vec}(\text{"king"}) - \text{vec}(\text{"man"}) + \text{vec}(\text{"woman"}) = \text{vec}(\text{"queen"})$
- Captures **semantic relationships** automatically.
- Unlike TF-IDF, embeddings **understand meaning**.



Example: Getting word embeddings with spaCy

Try it out yourself.. Visualize your embeddings



Embeddings in communication science

Applications:

- **Sentiment analysis:** Understanding audience reactions on social media (Rudkowsky et al., 2018)
- **Topic modeling:** cross-lingual topics
- **Recommender Systems:** Suggesting content based on user preferences (Loecherbach and Trilling, 2020)
- **Text similarity:** Measuring the similarity between texts (Brinberg and Ram, 2021)

Note: We will discuss this next week.

Using pretrained embeddings

Why use pretrained models?

- Trained on massive datasets (Google News, Wikipedia, etc.).
- Capture **rich linguistic structures**.
- Reduce training time and improve performance.

Popular choices:

- Word2Vec (Google News)
- GloVe (Common Crawl, Wikipedia)
- BERT (contextualized embeddings)
- fastText (subword information)

Why Use Embeddings?

- Capture **semantic meaning** of words.
- Handle **synonyms** and **related words** effectively.
- Work well in NLP applications: **text classification**, **clustering**, **sentiment analysis**.
- Used in communication science for **media analysis**, **misinformation detection**, and **social network studies**.

Key takeaways

- Traditional methods (TF-IDF) are **limited in capturing meaning**.
- Word embeddings create **dense vectors** that capture relationships.
- Pretrained models like **Word2Vec, GloVe, and BERT** help analyze text effectively.
- Embeddings are **widely used** in NLP and communication science.



Thank you!!

Thank you for your attention!

- Questions? Comments?

References i

References

-  Brinberg, M., & Ram, N. (2021). **Do new romantic couples use more similar language over time? Evidence from intensive longitudinal text messages.** *Journal of Communication*, 71(3), 454–477.
<https://doi.org/10.1093/joc/jqab012>
-  Loecherbach, F., & Trilling, D. (2020). **3bij3 – Developing a framework for researching recommender systems and their effects.** *Computational Communication Research*, 2(1), 53–79. <https://doi.org/10.5117/ccr2020.1.003.loec>

References ii



Rudkowsky, E., Haselmayer, M., Wastian, M., Jenny, M.,
Emrich, Š., & Sedlmair, M. (2018). **More than bags of
words: Sentiment analysis with word embeddings.**
Communication Methods and Measures, 12(2-3), 140–157.