

# Computational Communication Science 2

## Week 3 - Lecture

### » (Soft) cosine and recommender systems«

---

Anne Kroon

a.c.kroon@uva.nl, @annekroon

April 14, 2025

Digital Society Minor, University of Amsterdam

# Today

Recap: Week 1+2

Cosine similarity

RecSys

Knowledge-based Recommender Systems

Content-based Recommender Systems

Wrap up



*Everything clear from last weeks?*

## Recap: Week 1+2

---

# What We Covered in the Last Weeks

## Text Preprocessing

- **Tokenization:** Splitting text into words or phrases.
- **Stopword Removal:** Filtering out common, low-information words.
- **N-grams:** Creating combinations of words (e.g., unigrams, bigrams, trigrams, collocations).
- **Stemming/Lemmatization:** Reducing words to their root or base form.

## Text representations (Numerical form)

- **Count Vectorization:** Representing text as word occurrence counts.
- **TF-IDF Vectorization:** Weighs word importance relative to the document and corpus.
- **Embeddings:** Vector representations capturing semantic meaning (spaCy).

# What We Covered in the Last Weeks

## Top down and bottom up approaches

- We have discussed top-down and bottom-up approaches....
- ...today we'll talk about a *bottom-up* approach:
  - cosine similarity
- ...second half of the course will focus on *top-down*
  - machine learning (vectorizers + pre-processing remain important!)

# The Bigger Picture

## Comparing CCS-1 and CCS-2

- CCS-1 focused on the *basics of programming* — like learning how to drive.
- CCS-2 focuses on *applying computational techniques to analyze textual data* — more like learning how the engine works.
- This course emphasizes conceptual understanding and methodological application.
  - You'll be tested on your ability to apply concepts — not on writing complex code from scratch.

# The Bigger Picture

## Comparing CCS-1 and CCS-2

- **CCS-2** introduces more abstract concepts. Yes, it can be challenging — but you **can** do hard things! Mastering these ideas is empowering.
- Topics like *vectorization*, *cosine similarity*, and *machine learning* aren't just for this course — they're valuable, in-demand skills that look great on your resume.
- If you're having trouble, please reach out — we're here to support you!
- **Have feedback or want a consultation?** Tell us how we can support you better:  
<https://forms.office.com/e/7fzzTvKGyT?origin=lpLink>



# Recap: Test Your Understanding of Vectors

## Did You Get It?

- Participate here:

<https://app.wooclap.com/JGMBTB?from=event-page>

Join this Wooclap event



1

Go to **wooclap.com**

2

Enter the event code in the top banner

Event code  
**JGMBTB**



1

Send @JGMBTB to  
0970 1420 2908

2

You can participate

Disable answers by SMS

# Document-Term Matrix Comparison (Count, TF-IDF, Embedding)

Count Vectorizer

Doc	cat	sat	dog
D1	1	1	0
D2	1	0	1
D3	0	1	1

Raw word frequencies

TF-IDF Vectorizer

Doc	cat	sat	dog
D1	0.58	0.58	0.00
D2	0.48	0.00	0.66
D3	0.00	0.48	0.66

Weights rare/important words

Embedding (spaCy avg.)

Doc	dim1	dim2	dim3
D1	0.31	-0.04	0.88
D2	0.27	-0.12	0.95
D3	0.26	-0.01	0.83

Dense semantic vectors (truncated)

## Cosine similarity

---

# Why Should We Care?

## Cosine Similarity in Action

- Ever wonder how Spotify knows your next favorite song?
- Or why Netflix keeps recommending crime thrillers?
- Behind the scenes, they're comparing items as vectors.
- Cosine similarity tells us how “close” two items are in meaning or content.

## Key Idea

Two documents (or songs, users, products) are similar if their vector directions are close — even if their values (lengths) differ.

# Applications of Cosine Similarity

## Where It's Used

- In industry:
  - Search engines (e.g., ranking relevant results)
  - Recommendation systems (e.g., suggesting similar content)
- In academia:
  - *Linguistic alignment in communication* (Brinberg & Ram, 2021)
  - *Overlap between political speech and public opinion* (Hager & Hilbig, 2020)

# From Vectorization to Similarity

## Good news: You already know (almost) everything!

- You've worked with vectors
- You already know how to **vectorize** your data using `CountVectorizer` and `TfidfVectorizer`
- Now, you can use those vectors to calculate **cosine similarity** — it's just one more step!

## What This Means

Once your text is in vector form, you can compare:

- How similar two documents are
- Whether two users talk alike
- Which sentences match a query best

# Mathematical Representation

$$\text{similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

- Measures cosine of the angle between vectors.
- 0 (orthogonal, dissimilar), 1 (identical, similar).

# Interpreting Cosine Similarity Values

## What the Numbers Tell You

Cosine similarity values range from 0 to 1:

- 1.0 → Texts are *identical in direction* (highly similar).
- 0.8 – 1.0 → Very similar content or style.
- 0.5 – 0.8 → Some overlap in meaning, but not identical.
- 0.0 – 0.5 → Low similarity; texts likely on different topics.
- 0.0 → No overlap in direction (completely dissimilar).

## Tip

Cosine similarity is all about the *direction* of vectors — so it's great when document length doesn't matter, but shared emphasis (on certain words or ideas) does.



# Cosine Similarity in Python

You've already seen how to vectorize text — now let's add one more line to compute similarity!

```

1  from sklearn.feature_extraction.text import CountVectorizer
2  from sklearn.metrics.pairwise import cosine_similarity
3  import pandas as pd
4  # Sample documents
5  documents = [
6      "When I eat breakfast, I usually drink some tea",
7      "I like my tea with my breakfast",
8      "She likes cereal and coffee"
9  ]
10 # Vectorize the text (bag-of-words)
11 vec = CountVectorizer(stop_words='english') # OR use TfidfVectorizer()
12 count_matrix = vec.fit_transform(documents)
13 # Compute cosine similarity
14 cos_sim = cosine_similarity(count_matrix) --> Only this line is new :-)
15 # Display as DataFrame for better readability
16 print(pd.DataFrame(cos_sim))

```

# Beyond Cosine: Introducing Soft Cosine Similarity

## Limitations of Cosine using CountVectorizer and TfidfVectorizer.

- Only works with exact word matches (e.g. "car"  $\approx$  "automobile").
- Doesn't capture deeper **semantic relationships**.

## Soft Cosine Similarity

- Uses **word embeddings** (spaCy) to measure similarity even with different words.
- Captures **synonyms**, related terms, and contextual meaning (Sidorov et al., 2014).

# Cosine vs Soft Cosine: Side-by-Side Comparison

## Comparing Two Sentences

*Sentence A:* "I drove my car to work."

*Sentence B:* "I drove my automobile to work."

### Regular Cosine Similarity

- Method: CountVectorizer or TfidfVectorizer
- Only exact word matches
- **Score: 0.67**
- Misses that "car" and "automobile" are related

### Soft Cosine Similarity (Embeddings)

- Method: Word Embeddings (spaCy)
- Captures semantic relationships
- **Score: 0.95**
- Understands that "car"  $\approx$  "automobile"

## Key Takeaway

Soft cosine captures **meaning**, not just word overlap — more powerful for nuanced language tasks.

# Let's put this into practice!

Check out the walkthrough here:

[https://github.com/uva-cw-ccs2/2425s2/blob/main/week03/  
exercise-lecture/cosine\\_similarity\\_WALKTHROUGH.ipynb](https://github.com/uva-cw-ccs2/2425s2/blob/main/week03/exercise-lecture/cosine_similarity_WALKTHROUGH.ipynb)

RecSys

---

Congratulations! You now have everything you need to build a recommender system.



# Recommender Systems in Communication Science

## New Research Questions

1. **Political communication and journalism.** E.g., crafting personalized news diets. However, this may impact the diversity of news diets and democracy (Locherbach & Trilling, 2018; Möller et al., 2018)
2. **Organizational and corporate communication.** E.g., applications in hiring and recruitment.
3. **Persuasive communication.** E.g., recommendation algorithms for tailored health interventions (Kim et al., 2019)
4. **Entertainment communication.** E.g., movie recommenders.

# Recommender Systems

Types of recommender systems (Locherbach & Trilling, 2018; Möller et al., 2018; Wieland et al., 2021)

1. 'Basic' knowledge-base recommender systems
2. Content-based recommender systems
3. Collaborative recommenders (not part of this course)



# Knowledge-based RecSys

---

# Knowledge-based recommender system

## When to use?

- To overcome the **cold start problem**; when we do not have ratings of individual users.
- Simple model. It does not rely on user's explicit or implicit ratings, but on specific queries.
- Typical use case: Real-estate. Buying a house is, for most families, a rare/ single event.

Recap: Week 1+2  
○○○○○○○

Cosine similarity  
○○○○○○○○○

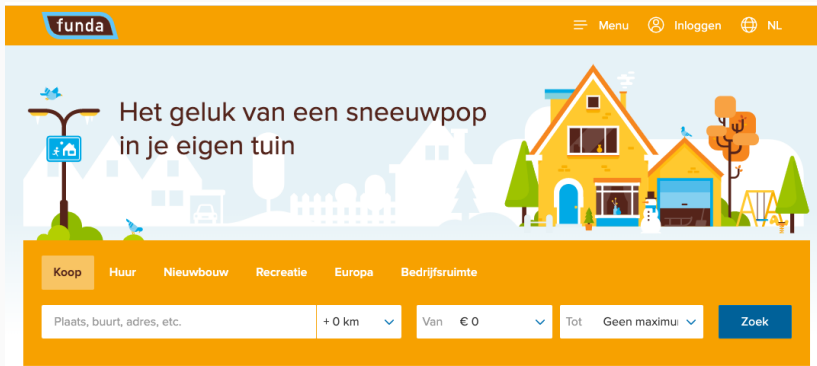
RecSys  
○○○○

Knowledge-based RecSys  
○○●○○○

Content-based RecSys  
○○○○  
○○○○

Wrap up  
○○○○

References



# Use case: IMDb database

	genres	title	tagline	release_date	vote_average	vote_count
0	[action, adventure, fantasy, science fiction]	Avatar	Enter the World of Pandora.	2009-12-10	7.2	11800
1	[adventure, fantasy, action]	Pirates of the Caribbean: At World's End	At the end of the world, the adventure begins.	2007-05-19	6.9	4500
2	[action, adventure, crime]	Spectre	A Plan No One Escapes	2015-10-26	6.3	4466
3	[action, crime, drama, thriller]	The Dark Knight Rises	The Legend Ends	2012-07-16	7.6	9106
4	[action, adventure, science fiction]	John Carter	Lost in our world, found in another.	2012-03-07	6.1	2124
...						



*What are relevant variables to use in a knowledge-based recommender system?*

# Knowledge-based recommender system

How can we work with user input without a front-end (such as the website of funda? → enter python's native `input()` function.

```
1 print("What is your favorite movie genre?")  
2 genre = input()
```

```
1 what is your favorite movie genre?  
2 [...]
```

# Improving knowledge-based recommender system

## When to use?

- It is important to think about ways to make the recommendation relevant for individuals
- Do you have more information in your db that make your top-listed recommendations as relevant as possible?

```
1 recommend_movies = movies.sort_values('vote_average',  
    ↪ ascending=False)
```

# Content-based RecSys

---





## Content-based systems

- Recommends items based on user's profiles.
- Profiles are based on e.g., ratings, and represents user's tastes/preferences.
  - For example, how often a user has clicked on, or liked, a movie.
- Recommendation is based on **similarity** between items in the content.
  - Content is here: e.g., genre, tags, plot, authors, directors, location, etc.

# Example of a content-based recsys

imdb.com/title/tt0241527/

Cast & crew · User reviews · Trivia · IMDbPro All topics



13 VIDEOS

99+ PHOTOS

Play trailer 0:32

Adventure Family Fantasy

An orphaned boy enrolls in a school of wizardry, where he learns the truth about himself, his family and the terrible evil that haunts the magical world.

**Director** Chris Columbus

**Writers** J.K. Rowling (novel) · Steve Kloves (screenplay)

**Stars** Daniel Radcliffe · Rupert Grint · Richard Harris

**IMDbPro** See production, box office & company info

Watch on Prime Video  
rent/buy from EUR2.99

+ Add to Watchlist

1.8K User reviews · 274 Critic reviews · 65 Metascore

# Example of a content-based recsys

## More like this



★ 7.5 ☆

Harry Potter and the  
Chamber of Secrets

Watch options



★ 7.9 ☆

Harry Potter and the  
Prisoner of Azkaban

Watch options



★ 7.7 ☆

Harry Potter and the  
Goblet of Fire

Watch options



★ 7.5 ☆

Harry Potter and the  
Order of the Phoenix

Watch options



# Content-based RecSys

---

Building blocks of content-based RecSys

## Feature selection and preprocessing

- Feature engineering is essential. What attributes or qualities do we want to include? In other words, which columns will you select and combine? (more on this tomorrow)
- Preprocessing your data (e.g., removing stop words, stemming) is important for improving cosine similarity with CountVectorizer and TF-IDF. However, for soft cosine (using embeddings), preprocessing is not required.

## Now how can we identify similar items?

1. Cosine similarity using text transformed with a CountVectorizer
2. Cosine similarity using text transformed with a TfidfVectorizer
3. Soft cosine similarity using word embeddings from an embedding model

## Benefits

- Content-based recommender systems are efficient and can provide highly personalized recommendations based on individual preferences.
- They are often integral to more complex recommender systems that combine deep learning and supervised learning techniques.

## You now know how to implement these!

- You've learned how to vectorize text and compute similarities — the foundation of content-based recommendations.
- With this knowledge, you can easily build effective, data-driven recommendation systems!

# Build your own recommender system

## Practice with the materials!

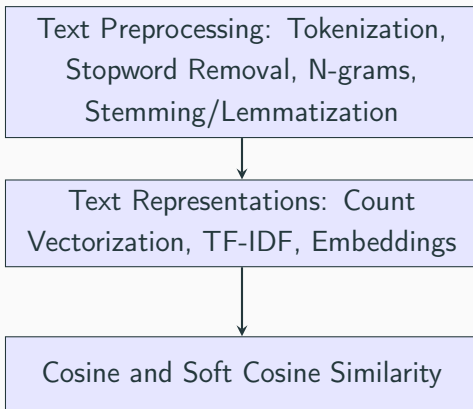
- To be able to do this correctly, it is essential that you understand the code of this week's lab session.
- Carefully walk through this week's assignment, and to whether questions arise.
- It's up to you to decide whether you want to build a simple knowledge-based or content-based recommender system. Base your selection on the available data columns.




Wrap up

---

## What We Covered in Part 1 – CCS 2: All ingredients of Content-based RecSys



## What's Next?



# Thank You!

Questions or feedback?

`a.c.kroon@uva.nl`

# References i

## References

---



Brinberg, M., & Ram, N. (2021). **Do new romantic couples use more similar language over time? Evidence from intensive longitudinal text messages.** *Journal of Communication*, 71(3), 454–477.

<https://doi.org/10.1093/joc/jqab012>



Hager, A., & Hilbig, H. (2020). **Does public opinion affect political speech?** *American Journal of Political Science*, 64(4), 921–937. <https://doi.org/10.1111/ajps.12516>

## References ii



Kim, H. S., Yang, S., Kim, M., Hemenway, B., Ungar, L., & Cappella, J. N. (2019). **An experimental study of recommendation algorithms for tailored health communication.** *Computational Communication Research*, 1(1), 103–129. <https://doi.org/10.5117/ccr2019.1.005.sukk>



Locherbach, F., & Trilling, D. (2018). **3bij3: A framework for testing effects of recommender systems on news exposure.** *Proceedings - IEEE 14th International Conference on eScience, e-Science 2018*, 350–351.  
<https://doi.org/10.1109/eScience.2018.00093>

## References iii



Möller, J., Trilling, D., Helberger, N., & van Es, B. (2018). **Do not blame it on the algorithm: an empirical assessment of multiple recommender systems and their impact on content diversity.** *Information Communication and Society*, 21(7), 959–977.

<https://doi.org/10.1080/1369118X.2018.1444076>



Sidorov, G., Gelbukh, A., Gómez-Adorno, H., & Pinto, D. (2014). **Soft similarity and soft cosine measure: Similarity of features in vector space model.** *Computacion y Sistemas*, 18(3), 491–504.

<https://doi.org/10.13053/CyS-18-3-2043>

## References iv



Wieland, M., Von Nordheim, G., & Kleinen-Von Königslöw, K. (2021). **One recommender fits all? An exploration of user satisfaction with text-based news recommender systems.** *Media and Communication*, 9(4), 208–221.  
<https://doi.org/10.17645/mac.v9i4.4241>